Introduction to Time Series Analysis and Data Types

Stochastic Process and Its Main Characteristics

A **stochastic process** is a collection of random variables indexed by time or space, representing a system that evolves over time in an uncertain manner. It is widely used in probability theory, statistics, and applications such as finance, physics, and engineering.

Main Characteristics of a Stochastic Process

1. **Index Set**: Represents the domain over which the process is defined. It can be discrete (e.g., time steps) or continuous (e.g., real numbers).
2. **State Space**: The set of all possible values the process can take. It can be discrete (e.g., a finite set of states) or continuous (e.g., real numbers).
3. **Stationarity**: A stochastic process is stationary if its statistical properties (e.g., mean and variance) do not change over time. It can be:
    - **Strictly Stationary**: The joint distribution of any subset of variables remains unchanged over time.
    - **Weakly Stationary**: The mean and autocovariance are time-invariant.
4. **Markov Property**: A stochastic process is Markovian if its future state depends only on the present state and not on past states.
5. **Ergodicity**: A process is ergodic if time averages and ensemble averages converge to the same value, meaning that statistical properties can be estimated from a single realization.
6. **Independent and Identically Distributed (i.i.d.)**: A process is i.i.d. if each random variable is independent and follows the same probability distribution.
7. **Memorylessness**: Some stochastic processes, such as Poisson processes and Markov chains, have the property that future states are independent of past states given the present state.

Time Series as a Discrete Stochastic Process

A **time series** is a sequence of observations recorded at discrete time intervals. It can be considered a **discrete-time stochastic process** where each observation is a realization of a random variable.

Types of Time Series

1. **Deterministic vs. Stochastic Time Series**
    - *Deterministic*: Values are predictable and follow a known pattern (e.g., a sine wave).
    - *Stochastic*: Values are influenced by random variations and uncertainty.
2. **Stationary vs. Non-Stationary Time Series**

- *Stationary*: The mean and variance remain constant over time.
- *Non-Stationary*: The statistical properties change over time (e.g., a trend or seasonal component).

3. **Examples of Stochastic Time Series Models**
    - **Autoregressive (AR) Model**: $Y_t = \phi Y_{t-1} + \epsilon_t$, where e_t is a white noise process.
    - **Example**:
    - An AR(1) process $Y_t = 0.8Y_{t-1} + \epsilon_t$,
      exhibits short-term dependence, decaying ACF, and weak stationarity if $|\phi|<1$ In contrast, a random walk has a variance that grows with tt, making it non-stationary.
    - **Moving Average (MA) Model**: $Y_t = \theta\epsilon_{t-1} + \epsilon_t$
    - **Autoregressive Moving Average (ARMA) Model**: A combination of AR and MA models.
    - **Autoregressive Integrated Moving Average (ARIMA) Model**: Extends ARMA by adding a differencing step to handle non-stationarity.

Time series analysis often involves identifying underlying stochastic processes, modeling dependencies, and forecasting future values.

1. What is Time Series Analysis?

Time series analysis is a statistical technique used to analyze temporal data. It involves examining time-ordered data points to identify patterns, trends, and seasonal variations to make predictions or understand underlying processes.

2. Objectives of Time Series Analysis

- **Descriptive Analysis**: Summarizing data behavior over time.
- **Forecasting**: Predicting future values.
- **Modeling Relationships**: Understanding dependencies between variables.
- **Control and Monitoring**: Identifying anomalies or shifts in trends.

3. Applications of Time Series Analysis

Time series analysis has a wide range of applications across various domains:

- **Finance:** Stock price prediction, risk management, and algorithmic trading.
- **Economics:** GDP forecasting, unemployment rate analysis, and inflation prediction.

- **Weather Forecasting:** Predicting temperature, rainfall, and other meteorological factors.
- **Healthcare:** Monitoring patient vitals, predicting disease outbreaks, and analyzing medical data over time.
- **Retail:** Sales forecasting, inventory management, and customer demand analysis.
- **Energy:** Electricity load forecasting, renewable energy production analysis, and energy consumption patterns.

4. Types of Data

1. **Time Series Data**: Data collected at different time points in a sequential manner (e.g., daily stock prices, monthly rainfall).

   Example:

Let's consider the monthly average temperature data for a city over a year. This data will help analyze temperature trends and seasonal variations.

Here's an example of time series data:

Month	Average Temperature (°C)
January	5
February	7
March	10
April	15
May	20
June	25
July	30
August	28
September	24
October	18
November	12
December	6

2. **Cross-Sectional Data**: Data collected at a single point in time across multiple subjects (e.g., income levels of different households in a city at one time).

   Example:

Imagine you're studying the impact of education level on income. You collect data on individuals' education levels and their corresponding incomes in the year 2025. Each individual's data is collected simultaneously, representing a cross-section of the population at that particular time.

Here's a simple example table of cross-sectional data:

Individual	Education Level	Income (USD)
Person 1	High School	30,000
Person 2	Bachelor's	50,000
Person 3	Master's	70,000
Person 4	PhD	90,000

This table shows the relationship between education level and income for different individuals at the same point in time.

TABLE 1.1

A Cross-Sectional Data Set on Wages and Other Individual Characteristics

obsno	wage	educ	exper	female	married
1	3.10	11	2	1	0
2	3.24	12	22	1	1
3	3.00	11	2	0	0
4	6.00	8	44	0	1
5	5.30	12	7	0	1
⋮	⋮	⋮	⋮	⋮	⋮
525	11.56	16	5	0	1
526	3.50	14	5	1	0

Aspect	Time Series Data	Cross-Sectional Data
Definition	Data collected over time at consistent intervals from the same subjects.	Data collected at a single point in time from multiple subjects.
Purpose	To analyze trends, patterns, and changes over time.	To analyze and compare differences between subjects at a particular point in time.
Structure	Each subject has multiple observations over different time periods.	Each subject has a single observation at one point in time.
Example	Monthly average temperature of a city over a year.	Survey of incomes of different individuals in a city in the year 2025.
Analysis	Allows for studying dynamics and trends over time.	Focuses on comparing characteristics between subjects at one time point.
Benefits	Identifies trends, patterns, seasonal effects, and long-term changes.	Provides a snapshot of data for comparison across subjects.
Complexity	Requires analysis of temporal patterns and potential autocorrelation.	Simpler analysis methods compared to time series data.
Visualization	Often visualized using line graphs to show changes over time.	Often visualized using bar charts, pie charts, or histograms to compare subjects.

3. **Pooled Data**: Combination of cross-sectional and time series data where repeated observations are taken over different subjects but not necessarily in a panel structure.

Pooled Data combines cross-sectional data from different time periods, so the same individuals or entities might appear, but there's no effort to track their individual changes over time.

Let's tweak the previous example for pooled data:

Year	Individual	Income (USD)
2020	Person D	45,000
2020	Person E	55,000
2020	Person F	50,000
2021	Person G	46,000
2021	Person H	56,000
2021	Person I	51,000
2022	Person J	47,000
2022	Person K	57,000
2022	Person L	52,000

Here, we see different individuals for each year. Pooled data aggregates different samples without following the same individual over time.

4. **Panel Data (Longitudinal Data)**: Data collected from the same subjects over multiple time periods (e.g., tracking economic growth of specific countries over years).

Panel Data captures multiple observations of the same subjects over time. This means each subject (like an individual) is followed and observed at several different time points.

Here's an updated example:

Year	Individual	Income (USD)
2020	Person A	45,000
2020	Person B	55,000
2020	Person C	50,000
2021	Person A	47,000
2021	Person B	57,000
2021	Person C	52,000
2022	Person A	49,000
2022	Person B	59,000
2022	Person C	54,000
2023	Person A	51,000
2023	Person B	61,000
2023	Person C	56,000
2024	Person A	53,000

Aspect	Panel Data	Pooled Data
Definition	Data collected from the same subjects over multiple time periods.	Cross-sectional data from multiple time periods combined into a single dataset.
Structure	Each subject has multiple observations over time.	Each observation is independent, potentially from different subjects across time periods.
Tracking Changes	Tracks changes for the same subjects over time.	Does not track individual changes over time.
Example	Annual income of the same individuals over five years.	Productivity data of different employees over several years, combined.
Analysis	Allows for studying individual dynamics and changes over time.	Focuses on aggregate analysis without individual tracking.
Sample Size	Limited to the same subjects over time.	Larger sample size by combining data from different subjects over time.
Complexity	Requires more complex statistical methods to account for time-series and cross-sectional variations.	Simpler analysis methods compared to panel data, but less detailed.

5. **Definition of Time Series**

A time series is a sequence of data points collected or recorded at specific time intervals. These data points are typically ordered in time, and the analysis involves understanding the underlying structure to make predictions or inferences.

5. Components of Time Series Data

A time series can be decomposed into several components, each representing a different aspect of the data:

- **Trend (T):** The long-term movement in the data. It represents the underlying direction (upward, downward, or stationary) over a long period.
- **Seasonality (S):** The repeating short-term cycle in the data. It is a pattern that repeats at regular intervals, such as daily, monthly, or quarterly.

- **Cyclical (C):** The fluctuations that are not of a fixed period. These are often influenced by economic or other external factors and can last for several years.
- **Irregular (I):** The random or noise component that cannot be explained by the trend, seasonality, or cyclical components. It represents the unpredictable variations in the data.

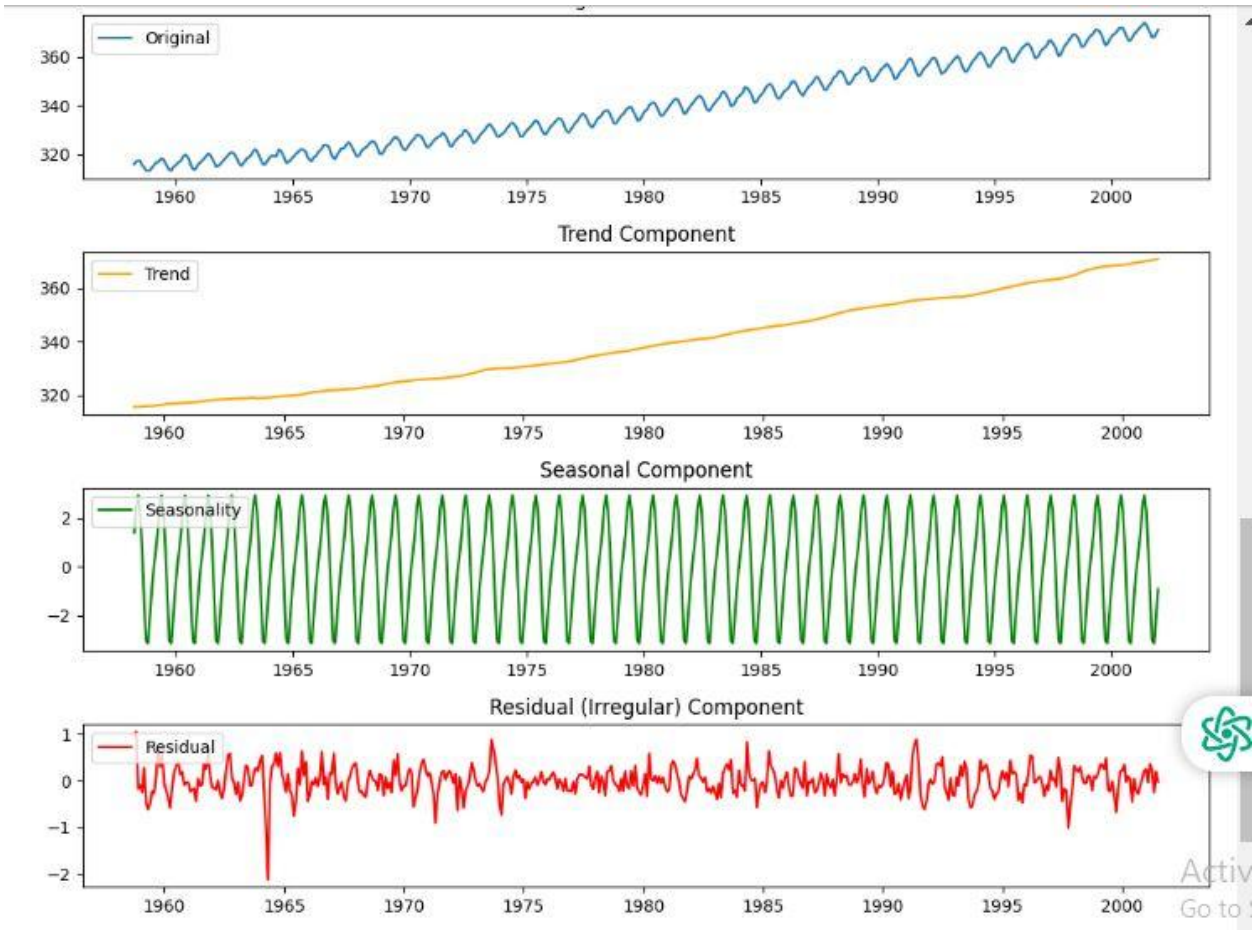

The decomposition can be additive or multiplicative:

- **Additive Model:** $Y(t)=T(t)+S(t)+C(t)+I(t)$
- **Multiplicative Model:** $Y(t)=T(t)\times S(t)\times C(t)\times I(t)$

Time series data can be decomposed into several key components:

1. **Trend (Tt)**: The overall direction of the data.
2. **Seasonality (St)**: Repetitive and predictable fluctuations.
3. **Cyclical Component (Ct)**: Long-term oscillations due to economic conditions.

4. **Irregular Component (It)**: Unpredictable, random variations.

Challenges in Time Series Analysis

- **Non-stationarity:** Many time series are non-stationary, meaning their statistical properties change over time. This requires transformations like differencing to make the series stationary.
- **Missing Data:** Handling missing data points is a common challenge, requiring imputation techniques.
- **Outliers:** Detecting and managing outliers that can skew the analysis.
- **High Dimensionality:** Dealing with multivariate time series where multiple variables are recorded over time.

5. Time Series Models

a) Additive Model

$Y_t = T_t + S_t + C_t + I_t$ Used when seasonal variations remain constant over time.

b) Multiplicative Model

$Y_t = T_t \times S_t \times C_t \times I_t$ Used when seasonal variations change proportionally with level.

Lecture on Time Series Analysis: Stationarity, ACF, and PACF

Introduction to Stationarity in Time Series

In time series analysis, stationarity is a fundamental concept that ensures the statistical properties of the series remain constant over time. A stationary time series is crucial for modeling and forecasting since many statistical models assume stationarity.

Definition of Stationarity

A time series Y_t is said to be stationary if its statistical properties, such as mean, variance, and autocorrelation, remain constant over time. Mathematically, a time series is stationary if:

• Constant Mean: $E(Y_t) = \mu$, for all t.

• Constant Variance: Var(Y_t)= σ^2 , for all t.

• Autocovariance depends only on lag: $Cov(Y_t, Y_{t+h})$= $\gamma(h)$which depends only on h(the lag) and not on actual t.

Types of Stationarity

• Strict Stationarity: The entire probability distribution of Y_t remains unchanged over time. This is a strong assumption and is rarely used in practice.
• Weak (Second-Order) Stationarity: The mean, variance, and autocovariance structure remain constant over time. This is commonly used in statistical modeling.

Example of a Stationary Time Series Theoretically

- Consider the white noise process:
- $Y_t = \in_t$, where $e_t \sim N(o, \sigma^2)$
  Here, the conditions for stationarity hold:
  $E(Y_t) = o$ (constant mean)
- $Var(Y_t) = \sigma^2$, (conatant variance)
- $Cov(Y_t, Y_{t+h}) = 0$, (zero autocovariance)

Non-Stationary Time Series

A time series is non-stationary if at least one of the stationarity conditions is violated. Common types of non-stationarity include:

• Trend Stationarity: The mean is not constant and changes over time (e.g., a linear trend).

• Seasonality: The series exhibits periodic fluctuations.

• Heteroskedasticity: The variance changes over time.

• Unit Root Processes: If a time series has a unit root, its mean and variance are undefined over time.

Example of a Non-Stationary Time Series

A random walk process:

$Y_t = Y_{t-1} + \in_t$
is non-stationary because its mean and variance increase over time.

The key properties of a random walk model include:

- Non-stationarity – The mean and variance are not constant over time.
- Expected Value (Mean)
- Variance
- Autocovariance
- Autocorrelation Function (ACF)

How to convert nonstationary series to stationary?

Differencing:

$\Delta Y_t = Y_t - Y_{t-1} = e_t$

Dtrending: Subtracting the trend component

Log transformation: To stabilize Variance

Simple Time Series Models

Time series models are used to describe and forecast data observed over time. Below are some fundamental models:

---

1. White Noise (WN) Model

A time series $\{Y_t\}$ is a **white noise process** if:

$$Y_t = \epsilon_t, \quad \epsilon_t \sim \text{i.i.d. } N(0, \sigma^2)$$

- **Characteristics:**
  - Mean: $\mathbb{E}[Y_t] = 0$
  - Constant variance: $\text{Var}(Y_t) = \sigma^2$
  - No correlation between values at different times.

2. Random Walk (RW) Model

A **random walk** is a simple non-stationary process where the current value depends on the previous value plus a random shock:

$$Y_t = Y_{t-1} + \epsilon_t, \quad \epsilon_t \sim \text{i.i.d. } N(0, \sigma^2)$$

- **Characteristics:**
  - Non-stationary (variance increases over time).
  - Common in financial markets (e.g., stock prices).
  - First difference ($\nabla Y_t = Y_t - Y_{t-1}$) is white noise.

3. Autoregressive (AR) Model

An **AR(p) model** expresses Y_t as a function of past values:

$$Y_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \cdots + \phi_p Y_{t-p} + \epsilon_t$$

For **AR(1)** (first-order autoregressive model):

$$Y_t = \phi Y_{t-1} + \epsilon_t, \quad |\phi| < 1$$

- **Characteristics:**
  - Stationary if $|\phi| < 1$.
  - Mean: $\mathbb{E}[Y_t] = 0$ (if no constant term).
  - Used for short-term dependencies.

4. Moving Average (MA) Model

An **MA(q) model** expresses Y_t as a function of past random shocks:

$$Y_t = \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \cdots + \theta_q \epsilon_{t-q}$$

For **MA(1)** (first-order moving average model):

$$Y_t = \epsilon_t + \theta_1 \epsilon_{t-1}$$

- **Characteristics:**
  - Short-term dependence structure.
  - Always stationary.
  - Used to model short-lived shocks.

5. Autoregressive Moving Average (ARMA) Model

A combination of **AR(p)** and **MA(q)**:

$$Y_t = \phi_1 Y_{t-1} + \cdots + \phi_p Y_{t-p} + \epsilon_t + \theta_1 \epsilon_{t-1} + \cdots + \theta_q \epsilon_{t-q}$$

- **Characteristics:**
  - More flexible than AR or MA alone.
  - Stationary if AR part is stationary.

6. Autoregressive Integrated Moving Average (ARIMA) Model

Handles non-stationary data by differencing d times:

$$\nabla^d Y_t = \phi_1 \nabla^d Y_{t-1} + \cdots + \phi_p \nabla^d Y_{t-p} + \epsilon_t + \theta_1 \epsilon_{t-1} + \cdots + \theta_q \epsilon_{t-q}$$

where $\nabla^d Y_t = (1 - B)^d Y_t$ represents differencing.

- **ARIMA(p,d,q):**
  - $p \rightarrow$ Autoregressive order
  - $d \rightarrow$ Differencing order
  - $q \rightarrow$ Moving average order
  - Used for non-stationary time series.

Proof of Non-stationary:

For a process to be stationary, it must satisfy three conditions:

- $E(Y_t)$ must be constant.
- $Var(Y_t)$ must be constant.
- The autocovariance function $\gamma(h) = Cov(Y_t, Y_{t+h})$ must depend only on the lag h, not on t.

We'll check these conditions for a random walk.

2.1. Expected Value (Mean)

Taking the expectation on both sides of the random walk equation:

$$E(Y_t) = E(Y_{t-1}) + E(\epsilon_t)$$

Since $E(\epsilon_t) = 0$, we obtain:

$$E(Y_t) = E(Y_{t-1})$$

Expanding recursively:

$$E(Y_t) = E(Y_{t-1}) = E(Y_{t-2}) = \ldots = E(Y_0)$$

Thus, the mean of the process depends on Y_0, which can change over time. Hence, the mean is **not constant**, violating stationarity.

Proof of Properties of White Noise (WN) Model

A white noise process $\{Y_t\}$ is defined as:

$$Y_t = \epsilon_t, \quad \epsilon_t \sim \text{i.i.d. } (0, \sigma^2)$$

where ϵ_t are independent and identically distributed (i.i.d.) random variables with mean 0 and variance σ^2.

Properties of White Noise and Their Proofs

1. Mean of White Noise

$$\mathbb{E}[Y_t] = 0$$

Proof:

Since $Y_t = \epsilon_t$ and $\epsilon_t \sim (0, \sigma^2)$, we have:

$$\mathbb{E}[Y_t] = \mathbb{E}[\epsilon_t] = 0$$

Thus, the mean of the white noise process is zero.

Variance of White Noise

$$\mathrm{Var}(Y_t) = \sigma^2$$

Proof:

By definition, the variance of Y_t is:

$$\mathrm{Var}(Y_t) = \mathbb{E}[Y_t^2] - (\mathbb{E}[Y_t])^2$$

Since $\mathbb{E}[Y_t] = 0$, we get:

$$\mathrm{Var}(Y_t) = \mathbb{E}[Y_t^2] = \mathbb{E}[\epsilon_t^2] = \sigma^2$$

Thus, the variance of the white noise process is **constant** and equal to σ^2.

. Autocovariance Function

$$\gamma(h) = \mathrm{Cov}(Y_t, Y_{t+h}) = \begin{cases} \sigma^2, & h = 0 \\ 0, & h \neq 0 \end{cases}$$

Proof:

The autocovariance function is given by:

$$\gamma(h) = \mathrm{Cov}(Y_t, Y_{t+h}) = \mathbb{E}[Y_t Y_{t+h}] - \mathbb{E}[Y_t]\mathbb{E}[Y_{t+h}]$$

Since $\mathbb{E}[Y_t] = 0$, we get:

$$\gamma(h) = \mathbb{E}[Y_t Y_{t+h}]$$

For $h = 0$:

$$\gamma(0) = \mathrm{Var}(Y_t) = \sigma^2$$

For $h \neq 0$: Since $Y_t = \epsilon_t$ are independent for all t,

$$\mathbb{E}[Y_t Y_{t+h}] = \mathbb{E}[\epsilon_t \epsilon_{t+h}]$$

Since ϵ_t and ϵ_{t+h} are independent,

$$\mathbb{E}[\epsilon_t \epsilon_{t+h}] = \mathbb{E}[\epsilon_t]\mathbb{E}[\epsilon_{t+h}] = 0 \cdot 0 = 0$$

Thus,

$$\gamma(h) = 0 \quad \text{for } h \neq 0$$

This proves that white noise has **no autocorrelation** except at lag $h = 0$.

4. Autocorrelation Function (ACF)

$$\rho(h) = \frac{\gamma(h)}{\gamma(0)} = \begin{cases} 1, & h = 0 \\ 0, & h \neq 0 \end{cases}$$

Proof:

The autocorrelation function (ACF) is defined as:

$$\rho(h) = \frac{\gamma(h)}{\gamma(0)}$$

From the previous result,

$$\rho(0) = \frac{\sigma^2}{\sigma^2} = 1$$

and for $h \neq 0$,

$$\rho(h) = \frac{0}{\sigma^2} = 0$$

Thus, the autocorrelation function of white noise is **1 at lag 0 and 0 otherwise**, meaning white noise has **no dependence over time.**

The white noise process Y_t satisfies the following properties:

1. **Zero Mean:** $\mathbb{E}[Y_t] = 0$.
2. **Constant Variance:** $\mathrm{Var}(Y_t) = \sigma^2$.
3. **Zero Autocovariance:** $\gamma(h) = 0$ for $h \neq 0$.
4. **No Autocorrelation:** $\rho(h) = 0$ for $h \neq 0$.

These properties make white noise a fundamental building block in time series analysis

Lag Operator in Time Series Analysis

The **lag operator** (also called the **backshift operator**) is a fundamental tool in time series analysis, used to simplify notation and manipulate time-dependent relationships.

Definition

The **lag operator**, denoted as B, shifts a time series backward by one period:

$$BY_t = Y_{t-1}$$

For example, if we apply the lag operator twice:

$$B^2Y_t = BY_{t-1} = Y_{t-2}$$

More generally, applying the lag operator k times results in:

$$B^kY_t = Y_{t-k}$$

Properties of the Lag Operator

1. **Linearity**: The lag operator follows the rule:

   $$B(aY_t + bX_t) = aBY_t + bBX_t$$

   for any constants a, b.

2. **Powers of B**:

   $$B^kY_t = Y_{t-k}$$

   This means repeated applications of B move the series further back in time.

3. **Inverse of the Lag Operator**: If the process is invertible, we can define the inverse of B:

   $$B^{-1}Y_t = Y_{t+1}$$

   This moves the series forward in time.

Application in Time Series Models

1. Autoregressive (AR) Model using the Lag Operator

Consider the AR(1) model:

$$Y_t = \phi Y_{t-1} + \epsilon_t$$

Using the lag operator, we rewrite this as:

$$Y_t - \phi B Y_t = \epsilon_t$$

or

$$(1 - \phi B) Y_t = \epsilon_t$$

For a general AR(p) model:

$$Y_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \cdots + \phi_p Y_{t-p} + \epsilon_t$$

Using B, this becomes:

$$(1 - \phi_1 B - \phi_2 B^2 - \cdots - \phi_p B^p) Y_t = \epsilon_t$$

2. Moving Average (MA) Model using the Lag Operator

For an MA(1) model:

$$Y_t = \epsilon_t + \theta_1 \epsilon_{t-1}$$

Using the lag operator, this is:

$$Y_t = (1 + \theta_1 B)\epsilon_t$$

For an MA(q) model:

$$Y_t = (1 + \theta_1 B + \theta_2 B^2 + \cdots + \theta_q B^q)\epsilon_t$$

3. Autoregressive Moving Average (ARMA) Model using the Lag Operator

An ARMA(p, q) model is:

$$Y_t - \phi_1 Y_{t-1} - \cdots - \phi_p Y_{t-p} = \epsilon_t + \theta_1 \epsilon_{t-1} + \cdots + \theta_q \epsilon_{t-q}$$

Using B, this becomes:

$$(1 - \phi_1 B - \cdots - \phi_p B^p) Y_t = (1 + \theta_1 B + \cdots + \theta_q B^q)\epsilon_t$$

4. Differencing Operator using B (For ARIMA Models)

To convert a non-stationary time series into a stationary one, we use **differencing**, denoted as:

$$\nabla Y_t = Y_t - Y_{t-1}$$

Using the lag operator:

$$\nabla Y_t = (1 - B)Y_t$$

For second-order differencing:

$$\nabla^2 Y_t = (1 - B)^2 Y_t = Y_t - 2Y_{t-1} + Y_{t-2}$$

This is useful in **ARIMA(p,d,q) models**, where d represents the number of differencing steps required to achieve stationarity.

The **lag operator** B simplifies time series notation and makes it easier to manipulate models. It is particularly useful for expressing:

- **Autoregressive models:** $(1 - \phi B)Y_t = \epsilon_t$
- **Moving average models:** $Y_t = (1 + \theta B)\epsilon_t$
- **ARMA models:** $(1 - \phi B)Y_t = (1 + \theta B)\epsilon_t$
- **Differencing operations:** $\nabla Y_t = (1 - B)Y_t$

Variance:

rencing, detrending, or transformations.

References

- Hyndman, R.J., & Athanasopoulos, G. (2018). *Forecasting: principles and practice*. OTexts.
- Box, G.E.P., Jenkins, G.M., & Reinsel, G.C. (2008). *Time Series Analysis: Forecasting and Control*. Wiley.